



MICRO 2024

# Cambricon-C: Efficient 4-bit Matrix Unit via Primitivization

Yi Chen

USTC & ICT, CAS





MICRO 2024

# Cambricon-C: Efficient 4-bit Matrix Unit via Primitivization

Yi Chen  
*USTC SKLP, ICT, CAS*  
chen\_yi@mail.ustc.edu.cn

Yongwei Zhao  
*SKLP, ICT, CAS*  
zhaoyongwei@ict.ac.cn

Yifan Hao  
*SKLP, ICT, CAS*  
haoyifan@ict.ac.cn

Yuanbo Wen  
*SKLP, ICT, CAS*  
wenyuanbo@ict.ac.cn

Yuntao Dai  
*USTC*  
daiyuntao@mail.ustc.edu.cn

Xiaqing Li  
*SKLP, ICT, CAS*  
lixiaqing@ict.ac.cn

Yang Liu  
*SKLP, ICT, CAS UCAS*  
liuyang22z1@ict.ac.cn

Rui Zhang  
*SKLP, ICT, CAS*  
zhangrui@ict.ac.cn

Mo Zou  
*SKLP, ICT, CAS*  
zoumo@ict.ac.cn

Xinkai Song  
*SKLP, ICT, CAS*  
songxinkai@ict.ac.cn

Xing Hu  
*SKLP, ICT, CAS SHIC*  
huxing@ict.ac.cn

Zidong Du  
*SKLP, ICT, CAS SHIC*  
duzidong@ict.ac.cn

Huaping Chen  
*USTC*  
hpchen@ustc.edu.cn

Qi Guo\*  
*SKLP, ICT, CAS*  
guoqi@ict.ac.cn

Tianshi Chen  
*Cambricon Technologies*  
tchen@cambricon.com





MICRO 2024

# Outlines

---


- Background & Motivation
- Algorithm
  - Primitivized Matrix Multiplication
  - Quarter Square Multiplication
- Architecture
  - Overview
  - PE design
  - Converter
- Evaluation and Results
- Conclusion

# Background: Emerging 4-bit Large Model



## Model Size Increases

- Model performance is related to model scale
- Model size imposes significant demands on the computational resources





## Model Deployment

- 4-bit variants get best performance under the same memory budget.
- GMML, vLLM, TensorRT-LLM, Open-VINO



# Background: Hardware Matrix Unit



## 1. Repeating Computation

- INT4 multiplication presents < 256 unique value.
- Matrix dimensions are often far > 256



**27.5% Tax !**

## 2. Accumulator dominating

- $n$  dimension INT4 matrix needs  $8 + \log_2 n$  bits Accumulator



# Algorithm: Primitivized Matrix Multiplication

MICRO 2024

## Key1: Avoid duplicated multiplications by pre-computing a Look-up Table

Suppose an INT4 Multiplication:

$$y = \sum_{i=0}^{N-1} w_i x_i, \quad N \gg 256$$

Define pre-computed  $\mathbf{m}$  includes all  $w_i x_i$  possible values

$$\mathbf{m} = [ \begin{array}{cccccc} 0 \times 0, & 0 \times 1, & 0 \times 2, & \dots, & 0 \times -1, \\ 1 \times 0, & 1 \times 1, & 1 \times 2, & \dots, & 1 \times -1, \\ 2 \times 0, & 2 \times 1, & 2 \times 2, & \dots, & 2 \times -1, \\ \vdots & \vdots & \vdots & \vdots & \vdots \\ -1 \times 0, & -1 \times 1, & -1 \times 2, & \dots, & -1 \times -1 \end{array} ]$$



Formula can also be presented as:

$$y = \sum_{i=0}^{N-1} \mathbf{p}_i \mathbf{m}^T,$$

Where  $\mathbf{p}_i$  is the one-hot vector


$$\mathbf{p}_i = [p_0, p_1, \dots, p_{255}], p_j = \begin{cases} 1, & \text{if } j = x_i \circ w_i \\ 0, & \text{otherwise} \end{cases}$$

## Key2: Reduce the strength of additions

With the distributive law of multiplication:

$$y = \sum_{i=0}^{N-1} \mathbf{p}_i \mathbf{m}^T = (\sum_{i=0}^{N-1} \mathbf{p}_i) \mathbf{m}^T$$

Major computation part: summation of **one-hot** vector  
 $\Rightarrow$  256 number increment + final once multiplication




(\* required in the last step multiplying  $\mathbf{m}^T$ , 8b  $\times$  16b to be precise.)



MICRO 2024



# Algorithm: Primitivized Matrix Multiplication





MICRO 2024


# Algorithm: Primitivized Matrix Multiplication






MICRO 2024

# Algorithm: Primitivized Matrix Multiplication




Compute  $\sum_{i=0}^{N-1} p_i$





MICRO 2024



# Algorithm: Primitivized Matrix Multiplication





MICRO 2024



# Algorithm: Primitivized Matrix Multiplication





# Algorithm: Primitivized Matrix Multiplication

MICRO 2024



# Algorithm: Quarter Square Multiplication

## Key: Reduce the number of counters

It is intuitively obvious that

$$\frac{(x+y)^2}{4} - \frac{(x-y)^2}{4} = xy$$

As  $x$  and  $y$  are integers.

$\Rightarrow x + y$  and  $x - y$  are both odd or even

$\Rightarrow$  The remainders will be cancelled.

$$\text{i.e. } (x+y)^2 \equiv (x-y)^2 \pmod{4}$$


$$\Rightarrow \left\lfloor \frac{(x+y)^2}{4} \right\rfloor - \left\lfloor \frac{(x-y)^2}{4} \right\rfloor = xy, \quad x, y \in \mathbb{Z}$$

The pre-compute look-up table becomes  $Q(n) = \lfloor n^2/4 \rfloor$

As  $Q(0) = Q(1) = 0$  and  $Q(n) \equiv Q(-n)$


$x + y \in [-16, 14]$ ,  $x - y \in [-15, 15]$ ,

$$\Rightarrow Q = [1, 2, 4, 6, 9, 12, 16, 20, 25, 30, 36, 42, 49, 56, 64]$$






# Architecture: Overview



- **Cambricon-C ACC**
  - Cambricon-C SA
    - 32x32 PE array
    - Converter
  - 128KB NBIn
  - 128KB NBOut
  - 256KB SB
  - Controller (Ctrl)
  - SFU
  - 8KB Instruction buffer (IB)
  - Direct memory access module (DMA)




# Architecture: Counters




## SRAM-based Counters

- + Lower leakage power
- + Lower area
- Inefficient incrementing
- Hard to customize



## Ripple Counters

- + Efficient incrementing
- + High speed
- + Easy to customize
- Large area and leakage power
- Unstable while incrementing




## Bi-directional Counters

- + Combine up/down counters
- + Easy to customize
- High leakage power
- Less efficient.



# Architecture: PE Designs




PE Design	Area ( $\mu\text{m}^2$ )	Energy (pJ)	Frequency (MHz)
MAC	<b>1422</b>	2.508	357
S256	21480	3.367	<b>400</b>
R225	57113	2.525	<b>400</b>
S32	13045	2.439	<b>400</b>
R29	8914	<b>1.274</b>	<b>400</b>
B15	9821	2.734	300

- Sacrifice area to gain **lower energy**
- **QSM** significantly increases the PE performance.
- **R29** outperform other schemes
- **S32** and **B15** has larger area and power. As it introduces more combination circuit.



# Architecture: Converter



- Converter for R29-PE array
  - Counter Readout
  - Subtraction of counter values
  - Multiplication by Pre-loading Value
  - Final summation



# Evaluation: Implementation


## Software

Model	# Param.	Quantization	Avg. Rank
VGG-16	138 M	LSQ* [11]	1141
ResNet-18	11.2 M		730
BERT-Base	110 M	KDLSQ* [22]	789
BERT-Large	340 M		1045
OPT-1.3B	1.3 B		2235
OPT-2.7B	2.7 B	QLLM [27]	2664
LLaMA2-7B	6.6 B		4013
LLaMA2-13B	12.7 B		5002

- 8 DNN models
- Compute same output value with the same quantization scale factor. (Increment will not be interrupted)
- **Rank:** number of MAC operations to compute one output value in matrix multiplication.


## Hardware

- PE: MAC, SIP, S256, R225, S32, R29, and B15
- ACC: TPU, Stripes, Camb-C ACC (R29 Based)
- Implementation:
  - SRAM: Open-RAM and CACTI-7
  - Other designs: Verilog
  - Synthesis: Design Compiler
  - PDK: FreePDK 45nm
- Simulation:
  - Power: gate-level VCD and Prime-Time PX
  - Clock Frequency: SDF file and VCS
  - Simulator: DNNWeaver and ANT






# Evaluation: Matrix dimension



- R29 outperform other PEs when matrix dimensions reach 128.
- The proportion of conversion energy decreases with the dimension increasing.
- The maximum benefit of R29 reaches 1.95x than MAC and 1.76x than SIP

# Evaluation: Performance on DNN



- Cambicon-C consistently outperforms TPU and Stripes in various DNN models.
- Cambicon-C gains 1.15x and 1.25x improvement compared with Stripes and TPU.



# Evaluation: Comparison with Other Methods

Design	Energy (pJ)	Frequency (MHz)
MAC [33]	2.508	357
SIP [23]	2.265	400
PIP [1]	2.549	400
FuseKNA [40]	6.450	400
CARAT [34]	6.638	400
R29	<b>1.273</b>	400

- PIP: Skip zero-bit in bit-serial compute.
  - Extra decoder and shifter
- FuseKNA: compute with Look-up table
  - LUT-based AGU involves multiple registers
- CARAT: Avoid repeating value computation
  - Suitable for float point data, but not INT.
  - longer array size to shed through put loss.

The circuit design for INT4 Computation is quite simple and neat, making it not easy to achieve net gains.



MICRO 2024

# Conclusion

- Algorithm:
  - Primitivized matrix multiplication
  - Quarter square multiplication
- Architecture
  - 5 different PE design
  - Converter
  - Cambricon-C accelerator
- Evaluation and Results
  - R29 gets 1.95x energy saving than MAC
  - Cambricon-C ACC get 1.25x energy saving than TPU



MICRO 2024

---

# Thanks!

# Q & A

Feel free to contact me on  
Whova app or e-mail: [chen\\_yi@mail.ustc.edu.cn](mailto:chen_yi@mail.ustc.edu.cn)